

Springleaf Customer Marketing Response

Robert Ciesielski

Executive Summary:

The Springleaf company would like to improve its targeting efforts by focusing on the customers who are likely to respond to direct mail offers. We propose a model that allows Springleaf to target 30% of the customers of interest and expect positive responses for 70% of the offers sent. Our result ranks on the 13th position of the private leaderboard of the corresponding Kaggle competition.

Contents

1	Introduction	1
2	Dataset, software and computing	1
3	Data Wrangling	2
4	Exploratory Data Analysis and Feature Engineering	2
4.1	Categorical variables	2
4.2	Numerical variables	3
4.2.1	DateTime features	3
4.2.2	Floats	4
4.2.3	Integers	4
5	Dimensionality Reduction	5
5.1	Correlation Matrix Method	5
5.2	Principal Component Analysis (PCA)	5
5.3	Logistic Regression with L1 Regularization	5
5.4	Recursive Feature Elimination (RFE)	5
6	Modeling	6
6.1	Logistic Regression Classifier (Baseline)	6
6.2	Random Forest Classifier	6
6.3	Gradient Boosting Classifiers	7
6.3.1	XGBoost Model	7
6.3.2	LightGBM Model	7
6.4	Results	8
6.4.1	Confusion Matrix	8
6.4.2	Feature Importance	9
6.5	Neural-Network Based Classifier	9
7	Summary	9

1. Introduction

Springleaf Financial (currently OneMain Financial¹) is a company that provides personal loans and optional insurance products to customers with limited access to traditional lenders, such as banks or credit card companies. The Springleaf's main marketing strategy is based on direct offers sent to customers by mail. The company is interested in improving its targeting efforts by efficiently focusing on the customers who are likely to respond to these offers. At the end of 2015, the company released an anonymized dataset with about 2,000 features to the Kaggle

community², asking it to predict which customers will respond to a direct mail offer. The model evaluation metric suggested by Springleaf is the area under the receiver operator curve, auc_roc.

This report describes our effort in modeling customer adoption using the Springleaf dataset. We focus on finding an optimal method to reduce the dimensionality of this dauntingly wide dataset. We compare the performance of various ensemble-based models, such as Random Forest, two scikit-learn Gradient Boosting models, XGBoost and LightGBM. The customer response is also modeled using a feed-forward neural network architecture with three and four hidden layers.

2. Dataset, software and computing

The data is available from the Kaggle platform as two csv files, one for model training and the other for model prediction and evaluation. The training dataset contains more than 145,000 rows corresponding to customers and 1932 columns with features describing the customers. The features are anonymized and numbered as integers. An additional column describes the binary customer response for supervised learning. About 23 percent of the customers respond positively to the offer (Tab. 1).

The computation is done on a PC with 8 CPU units and 16 GB RAM, using latest versions of pandas, numpy, scipy, scikit-learn, xgboost, lighgbm, keras and tensorflow libraries, as implemented in the anaconda v1.7.2 package. The hyperparameter tuning is computed on the cloud using the SegMind platform³. An instance of 32 CPU units and 128 GB RAM allows us to increase performance speed by a factor of 5. When running the BayesSearchCV function from the latest scikit-optimize (0.8.1) package, the scikit-learn and scipy libraries are downgraded

²<https://www.kaggle.com/c/springleaf-marketing-response/overview>

³<https://segmind.com/>

Customer response	N. of rows
No (0)	111,458 (76.7%)
Yes (1)	33,773 (23.3%)

Table 1. Binary customer response (target variable).

¹https://en.wikipedia.org/wiki/OneMain_Financial




Figure 1. Number of feature unique values.

(from versions 0.24.1 and 1.5.3 to 0.23.2 and 1.4.1, respectively) to assure the compatibility with scikit-optimize.

3. Data Wrangling

We begin with examining the dataset in terms of missing values, feature types, and unique values per feature. Fig.1 shows a histogram with the number of feature unique values. The spectrum is very wide and varies from 5 columns with just one unique value to a column with more than 100 thousand values. Features with only one unique value have no variance and are removed from the dataset. In this initial step, we inspect the content of the features that have at most five unique values. We find out that the majority of the two-value columns contain a NaN (not-a-number) and another value appearing as a string, a bool, or a number. For these columns, we preserve the NaNs but convert all the other values to 1. There are also four columns with the values of -99999 and 0. To unify the notation with the other columns, we replace -999999 by 1. A few columns contain logical variables denoted with boolean values False or True. We convert them to 0 and 1, respectively. Finally, we check for duplicates. After the replacement mentioned above, most of the columns with two unique values are duplicated, so we remove them from the dataset. We also find duplicate columns among those with 3 unique values. We remove them as well. None of the columns with 4 or 5 unique values are repeated. After this step, we are left with 1881 features (51 removed).

Tab. 2 shows the data type of the features in the dataset after the initial cleaning. Possible missing values are ignored when inferring the feature type. The dataset is dominated by integer columns (98%) with a small admixture of float and DateTime

Feature type		N. of features
Numerical:	Integers	1,837
	Floats	10
	DateTimes	18
Categorical:	Strings	17

Table 2. Feature data types.




Figure 2. Number of rows with missing values per column, with the missing values explicitly denoted as NaNs (top) , or represented by frequently repeated negative or positive numbers given in the legend (bottom). The dotted red lines correspond to the total number of rows in the datasets.

columns and categorical features represented by strings.

The number of missing values per feature is shown in Fig. 2 (top). Only the first 530 columns in the dataset contain missing values explicitly denoted as NaNs. Their number is very high for a few columns and close to the total number of dataset rows (red line). Most of the columns have the missing value fraction below 1%. Missing values could also be encoded as integers, e.g. as the -1 digit or a very small or large numbers ($1 - 10^n$ or $10^n - 1$ for $n = 2, 3, \dots$). By searching for the most frequently repeated column minimum and maximum values, we find those to be: -99999, -1 and 99, 999, 9999, 999999999. Fig. 2 (bottom) shows the frequency of their appearance. Overall, missing values are present in most of the columns in the dataset.

4. Exploratory Data Analysis and Feature Engineering

Given that the data is anonymized, we could blindly apply label- or one-hot encoding to convert categorical features into numerical ones and quickly proceed to the modeling stage. However, we will make an extra effort to explore and engineer the features, and hopefully get insight into the origin of the data. We inspect the columns by their types shown in Tab. 2 (in reversed order):

4.1 Categorical variables

Among 17 columns with string-encoded categorical features:

- One column contains more than 12,000 city names. Using an additional dataset from the US Geological Survey and




Figure 6. Left: the hour of the day versus the day of the month for entries in the column covering 3 days (notice the jump from Jan to Feb on the horizontal axis). Middle: the time correlation between two columns with the shortest time intervals, explained in the text. Right: Average target response versus the number of non-zero entries in the DateTime columns summed over the rows.

significant number of empty rows (70-99%). One of the fully-filled columns spans the time range of merely 3 days, another one of about a year, while all the others cover periods of 4 to 13 years. Fig. 6 (left) shows the hour of the day versus the day of the month for entries in the column that covers 3 days. There is a one-to-one correspondence between this column and the column that spans one year, as shown in the middle plot. No obvious correlations are found between any other columns. Under the assumption that during 3 afternoon shifts at the end of Jan 2014, somebody compiled information about other chronologically ordered data, we remove the first column from the dataset. For other columns, we convert the DateTime entries to floats by taking the difference w.r.t. to the smallest date within each column. Missing values are replaced by -1. Moreover, we count non-zero DateTime entries in each row and save the counts as an extra feature. Fig. 6 (right) shows a strong dependence of average target response on these counts. They may represent the number of treatments a customer underwent (the more treatments, the higher the probability of the positive customer response).

4.2.2 Floats

There are ten columns in this class and all exhibit steeply falling exponential distributions. One notable feature is that they contain two types of missing values: NaNs and -1 digits. The average target response for rows with these values differs from the overall average by 2 and 1%, respectively. For this reason, we replace NaNs with -2 to keep them separate from the -1 values.

4.2.3 Integers

This class is present in 98% of the columns. Since it would be impractical to inspect all of them in detail, we focus on some aggregated information. We group the columns by missing values that they contain, encoded as (see Sec. 3):

- negative numbers: -1 and -99999; Fig. 7 shows standard box plots (with the median, 1st and 3rd quantiles, Tukey's whiskers, etc.) for 104 and 156 columns containing these numbers, after removing them from the spectra. The columns are ordered with increasing median values, showing a very large variation in the range of distributions, filled with positive numbers. A visual

inspection of selected columns shows that they are dominated by steeply falling spectra, most likely representing medical measurements. There are also columns with numbers that look like, e.g., a year or a ZIP code. While the -1 digit seems adjacent to the distributions, the -99999 number lies five orders of magnitude apart from the spectra, in most cases creating an enormous gap that might be problematic for modeling. To avoid possible issues, we replace it with -1.

- large positive numbers: 99, 999, 9999 and 999999999; present in 458, 238, 159 and 257 columns, respectively. Here, we discover a much more complicated structure, as these numbers are accompanied by other adjacent large numbers. The columns consist of a set of numbers (994, 995, 996, 997, 998, 999) or (9994, 9995, 9996, 9997, 9998, 9999) or (999999994, 999999995, 999999996, 999999997, 999999998, 999999999), which are separated from the rest of the spectra by a gap not larger than two orders of magnitude (for most cases). Such a picture is difficult to understand without additional domain knowledge, so we leave the data unchanged.
- explicit NaN values; in a small number of columns that contain them, we replace the NaNs with the -1 digits.

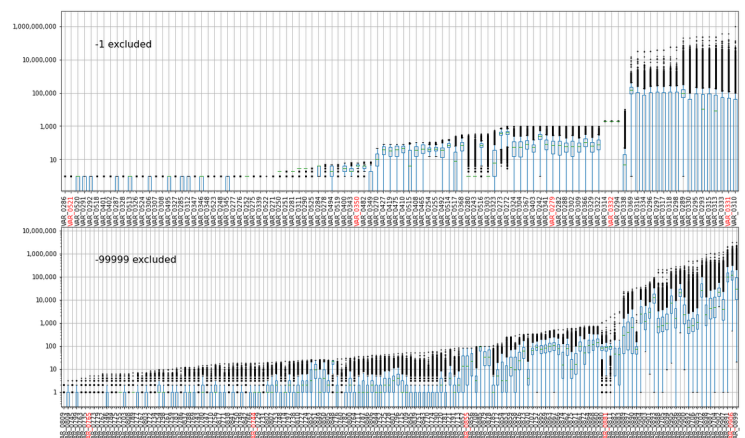


Figure 7. Box plots as a function of column ID for integer columns with missing values encoded as negative numbers: -1 (top) and -99999 (bottom), after excluding them from the data.

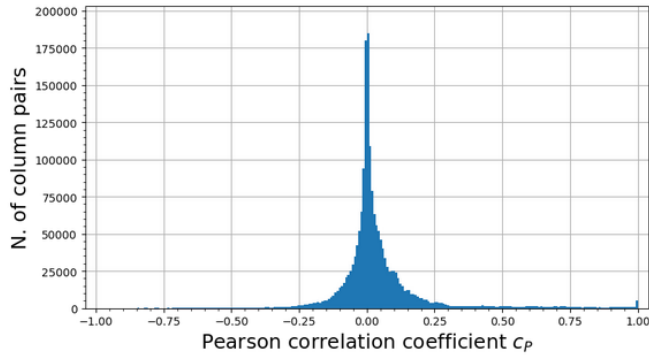


Figure 8. Pearson correlation coefficient for pairs of columns.

5. Dimensionality Reduction

The following four approaches are explored to reduce the dimensionality of the dataset, consisting of 1,892 columns.

5.1 Correlation Matrix Method

This method relies on a pair-wise comparison of columns in the dataset and the removal of highly similar ones. The level of similarity is measured in terms of the Pearson correlation coefficient c_P ($-1 < c_P < 1$), shown in Fig. 8. The distribution of c_P is centered at zero implying that the data consists of minimally correlated column pairs. A rather loose cut $|c_P| > 0.5$ would remove only 5% of columns, not sufficient for desired dimensionality reduction. In the plot, a small excess of columns at $c_P = 1$ corresponds to fully correlated pairs. We remove 12 duplicates from the dataset and are left with 1,880 columns.

5.2 Principal Component Analysis (PCA)

PCA allows us to create a new set of features (principal components, PCs) that are linear combinations of original features. Typically, fewer PCs than original features are needed to describe the variance of data fully. We exploit this property for dimensionality reduction. After scaling the data and applying PCA, we find that 500 and 700 PCs are needed to explain 97 and 99% of variance in the data, respectively. Unfortunately, we observe a significantly poorer modeling performance for the reduced datasets or even for the entire dataset after the transformation (all 1,880 PCs). For the latter, the auc_roc score of the Random




Figure 9. PC decomposition in the original feature basis (only the first 100 PCs are shown on the vertical axis).

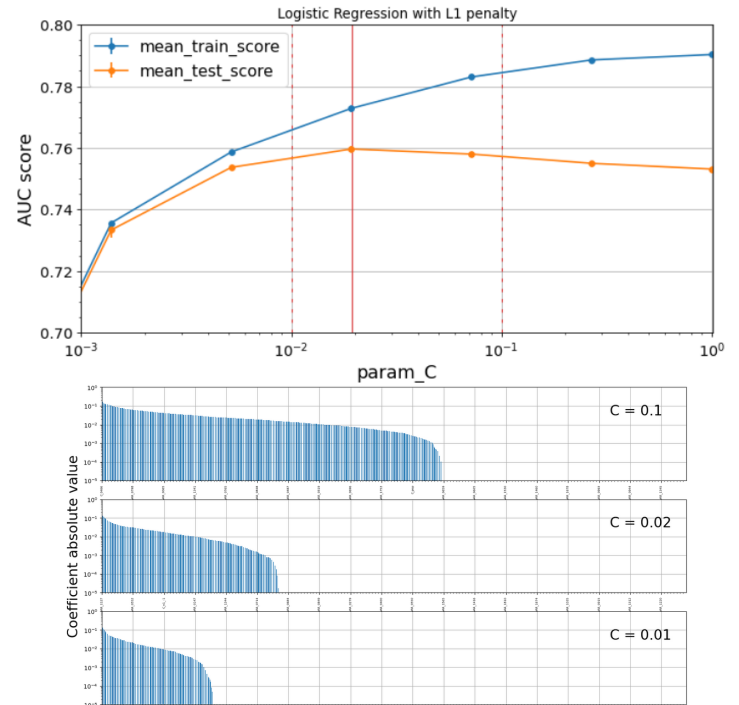


Figure 10. Score as a function of the L1 penalty parameter C (top) and feature coefficients (bottom) for the LR model.

Forest model drops from 0.76 to 0.67, while its run time increases by a factor of 5. We attribute this behavior to strong correlations between PCs, as shown in Fig. 9. Each PC is a combination of almost all the original features (non-zero coefficients on the z-axis), which affects the discriminative power of models. This could be improved by using a Sparse PCA method, but we decide to drop this approach and try other solutions.

5.3 Logistic Regression with L1 Regularization

A regularization by adding the L1 penalty to the Logistic Regression (LR) model offers a way to prevent overfitting and remove redundant features. The magnitude of the penalty is controlled by the parameter C . Fig. 10 (top) shows the results of the fit to the data with the LR model, as a function of C . The blue (yellow) curve corresponds to the auc_roc score obtained for the train (test) set. In the test data, the maximum score of 0.767 is obtained for $C=0.02$ (the solid red line). The score varies very slowly around the maximum and drops to 0.765 for $C=0.01$ and 0.1 (dotted red lines). Fig. 10 (bottom) shows distributions of absolute values of feature coefficients for fits with $C=0.01$, 0.02, and 0.1 (from top to bottom). Entries with 0 correspond to features removed by the model. The number of non-zero features amounts to 1096, 569, and 359, respectively. We select the latter two for the modeling in the next step.

5.4 Recursive Feature Elimination (RFE)

RFE is an iterative procedure in which, at each step, a model is evaluated after removing n least important features, until

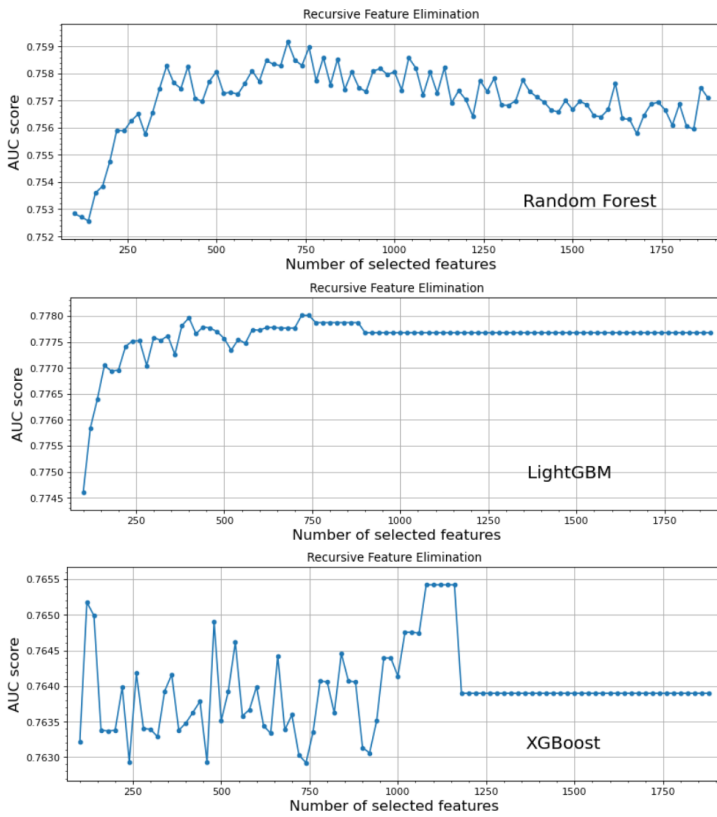


Figure 11. RFE with the Random Forest, LightGBM, and XGBoost models.

reaching a minimum number of features, n_{\min} . We apply RFE to three models: Random Forest, LightGBM, and XGBoost, with default values of hyperparameters. The step parameter n is set to 20, and n_{\min} to 100. Models are evaluated using the `auc_roc` metric and the two-fold cross-validation technique. The results are presented in Fig. 11, which shows that the maximal score is obtained when the number of selected features is 700, 360, and 1080 for Random Forest, LightGBM, and XGBoost models, respectively. For LightGBM, two regions obtain the maximal score. To access the one with a smaller number of features, we repeat the RFE procedure three times.

Tab. 3 summarizes methods used to reduce the dimensionality of the dataset. Dots indicate those that will be used for modeling (less than 1,000 features). Notice that out of 360 features selected by both LR, $L-0.01$ and RFE, LGBM models, only 150 (41%) are common.

Method	N. of columns
• (none)	1880
• LR, L1-0.1	1096
• LR, L1-0.02	569
• LR, L1-0.01	359
• RFE, XGB	1080
• RFE, RF	700
• RFE, LGBM	360

Table 3. Dimensionality reduction, summary.

6. Modeling

The customer response is modeled using decision-tree-based ensemble methods. We check the performance of the scikit-learn Random Forest (RF) classifier and four models based on the gradient boosting approach: two from the scikit-learn library, GradientBoosting (GB) and HistGradientBoosting (HGB), as well as XGBoost (XGB) and LightGBM (LGBM). The baseline model of our choice is Logistic Regression (LR). We also test a model based on the neural network (NN). Models are evaluated using the area-under-the-ROC-curve (`auc_roc`) score.

The hyperparameter tuning for ensemble models is conducted in three steps. First, we test the performance of the models with their default settings. Then, in a manual tune, one-dimensional parameter scans are done using the scikit-learn GridSearchCV method to select those parameters and their ranges that maximize the score function. Finally, a multi-dimensional Bayesian optimization of selected parameters is performed using the scikit-optimize BayesSearchCV method. Typically, the first step takes about a minute on our 8-CPU machine, the second one requires a few hours, while the third step would need five days (one day) of computing for the full (the smallest) dataset. For the latter case, we compute on the cloud (the Segmind platform), reducing the running time to one day (a few hours).

In the following, given a rather high complexity of the problem (several models, several datasets), we will first evaluate the performance of the above-mentioned models with the full dataset, then test the most promising model against the remaining reduced datasets.

6.1 Logistic Regression Classifier (Baseline)

Although the essential LR requirement of the linear dependence of the logit function (log-odds) on independent data features is most likely not met, the LR model provides a simple reference for ensemble models. It was introduced in Sec. 5.3. After scaling the data with the Standard Scaler and tuning the parameter C , the highest `auc_roc` score of 0.767 was achieved for $C=0.02$.

6.2 Random Forest Classifier

The `auc_roc` score for a single Decision Tree model is 0.582. The RF model with default parameters (100 trees) gives a score of 0.753. Fig. 12 (left) shows that the score increases with the additional number of trees until it reaches a plateau at about 1,000 trees. The plot is a part of the RF model's manual tune, which includes the seven most significant RF hyperparameters. During the tune, the score for each subsequent parameter is evaluated using the best values of the previous parameters⁵. The best tune yields the RF `auc_roc` score of 0.779 for $n_{\text{estimators}}=2000$, $\text{max_depth}=17$, $\text{max_features}=180$, and $\text{bootstrap}=\text{False}$ (all

⁵Of course, such an approach depends on the order of the parameters and does not ensure finding the score's global maximum. It can be thought of as a relatively quick local maximum search, sufficient for model comparison.




Figure 12. The auc_roc score as a function of the number of trees for RF and selected tree depths (left), and for XGBh and selected learning rates (right); entire dataset.

dataset rows considered in building a tree). These parameters differ significantly from their default values of 100, unlimited, $\sqrt{nf} = 43$, and bootstrap = True (a fraction of rows used), respectively. For other parameters, the default values work well. We do not perform the time-consuming Bayesian optimization here, as we expect other models to perform better than RF.

6.3 Gradient Boosting Classifiers

The four models under study are based on a similar principle but differ in detail. E.g., GB is the earliest implementation of the gradient boosting technique among the four. XGB simplifies the loss function minimization by using only two leading terms of its Taylor expansion, allowing more robust performance (scalability, parallelization, etc.). LGBM, rather than adding subsequent trees, expands one tree in depth with additional leaves and uses histogram representation of features for faster performance. HGB is an experimental version of a scikit-learn model inspired by LGBM, etc. The models also differ by parameter default values (and their names), as shown in Tab. 8 at the end of this report. To compare model performances, we set the default learning_rate and max_depth parameters to the common values of 0.1 and 5, respectively. By doing so, the number of leaves for GB and XGB is set to $2^{\text{max_depth}} = 32$, allowing a fair comparison with the LGBM and HGB models (default 31 leaves).

Tab. 4 shows the auc_roc, accuracy, precision, and recall scores, together with the computation time for each of the models. XGBh represents a variation of XGB with histogram representation of features, similarly to LGBM and HGB. Generally, the scores are pretty similar, with slightly better auc_roc results for XGBh and LGBM. Moreover, the GB model does not offer parallelization, so its computation time is at least eight times longer than other models. XGBh is about two times faster than

	auc_roc	acc.	prec.	recall	time (min.)
GB	0.777	0.794	0.628	0.281	18.2
XGB	0.777	0.795	0.633	0.279	1.9
XGBh	0.778	0.794	0.633	0.276	0.9
LGBM	0.778	0.795	0.629	0.287	0.4
HGB	0.777	0.795	0.629	0.285	1.0

Table 4. Scores for gradient boosting models using full dataset.

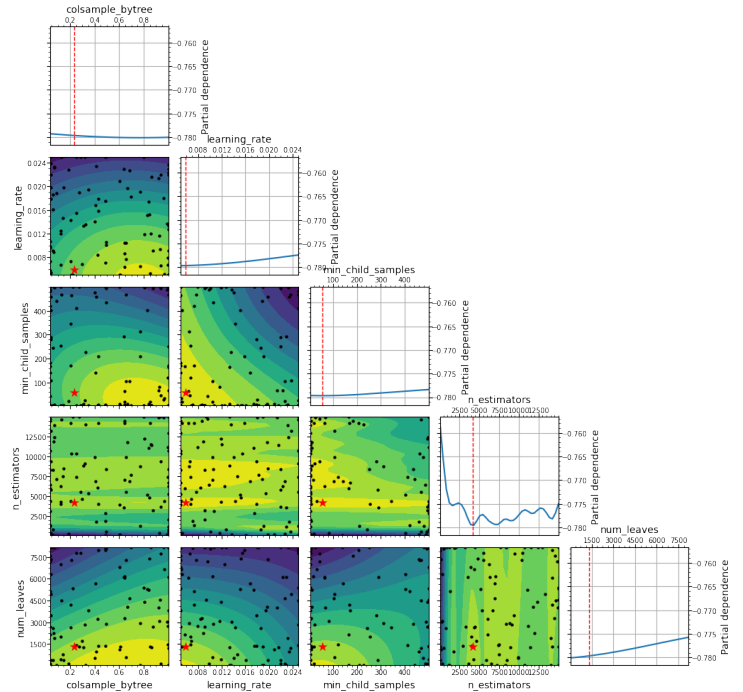


Figure 13. Partial dependence plots for Bayesian optimization of LGBM (full dataset): the auc_roc score as a function the shown parameter(s), averaged over the other parameters. Red stars and lines represent the best parameter values.

XGB, and LGBM about twice as fast as HGB. Given these observations, we discard the GB, HGB, and XGB models and in the following focus on XGBh and LGBM.

6.3.1 XGBoost Model

The auc_roc score for the default XGBh model obtained in the previous section is 0.7779. The manual tune, based on a scan of the eight most crucial parameters, increases the score to 0.7868. Fig. 12 (right) shows an example of the tune as a function of the number of trees for selected values of the parameter learning_rate. The Bayesian optimization yields the score of 0.7969 for learning_rate = 0.005, n_estimators = 3136, max_depth = 47, colsample_bytree = 0.27, min_child_weight = 5, reg_lambda = reg_alpha = 0, and other parameters set to their default values.

6.3.2 LightGBM Model

The auc_roc score for the default LGBM model is 0.7783. The manual tune (12 parameters) increases the score to 0.7884. The score after the Bayesian optimization is 0.7985, for learning_rate = 0.006, n_estimators = 4215, num_leaves = 1319, colsample_bytree = 0.23, min_child_samples = 57 and other parameters set to their default values. Fig. 13 shows one- and two-dimensional partial dependence plots for the five parameters after the fit, which converges after 50 iterations.

Hyperparameter set	RF	XGBh	LGBM
	score ratio	score ratio	score ratio
Default (D)	0.753 1.00	0.7779 1.00	0.7783 1.00
Manual Tune (MT)	0.779 1.03	0.7868 1.01	0.7884 1.01
Bayesian Opt. (BO)		0.7969 1.02	0.7985 1.03
D	0.97	1.00	1.001
MT, ratio to XGBh	0.99	1.00	1.002
BO		1.00	1.002

Table 5. Summary of the parameter optimization for the RF, XGBh and LGBM models (full dataset).

6.4 Results

Tab. 5 summarizes the results obtained so far for the RF, XGBh, and LGBM models. In the top rows, the auc_roc score is shown for the three different parameter sets, together with the ratio relative to the default set (D). The bottom rows show the score ratio between the models relative to XGBh. It is interesting to notice that a relatively quick manual tune improves the results by as much as 3% for RF and 1% for gradient boosting models. The time-involving Bayesian optimization yields an additional 1-2% for XGBh and LGBM. Moreover, RF performs worse than the gradient boosting models by a few percent (1-3%), while LGBM systematically outperforms XGBh by a few permille (1-2‰). We choose the LGBM model to study the impact of the dataset reduction on the customer response.

Tab. 6 presents the results of the Bayesian optimization with the LGBM model performed using the reduced datasets mentioned in Sec. 5. Five parameters are sampled in the common ranges of $0.005 < \text{learning_rate} < 0.025$, $100 < \text{n_estimators} < 15,000$, $16 < \text{num_leaves} < 2^{13} = 8192$, $5 < \text{min_child_samples} < 500$, and $f_{\min} < \text{colsample_bytree} < 1$. For the last parameter, the fraction f_{\min} is set such that for each dataset there are at least 70 features considered at each tree node. The best score is obtained for the reduced dataset with 360 features, extracted using the *RFE*, *LGBM* method. The corresponding auc_roc value is 0.7998, which is 2 permille higher than the score derived for the full dataset. It is obtained for learning_rate=0.0051, n_estimators=10,046, num_leaves=7,583, min_child_samples=30, and colsample_bytree=0.23, and represents the final result of this study. The result ranks on the

Dataset	(n. cols)	auc_roc	ratio to Full
Full	(1880)	0.7985	1.000
LR, L1-0.02	(569)	0.7962	0.997
LR, L1-0.01	(359)	0.7959	0.997
RFE, RF	(700)	0.7971	0.998
RFE, LGBM	(360)	<u>0.7998</u>	1.002

Table 6. LGBM after Bayesian optimization for reduced datasets defined in Tab. 3.

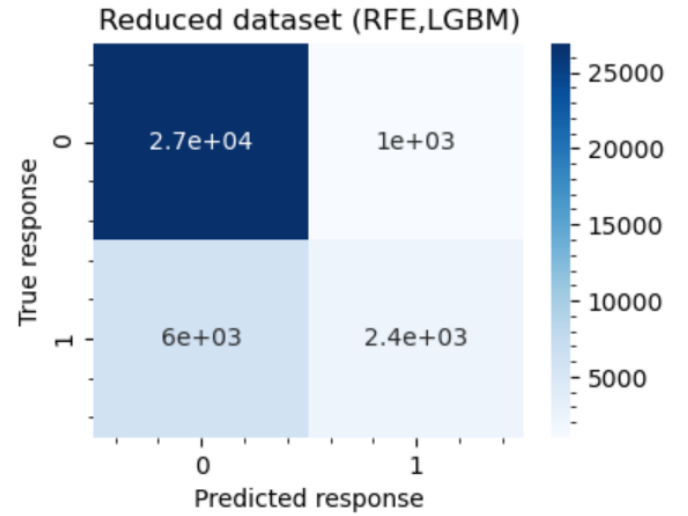


Figure 14. Confusion matrix, final model, test sample (36.3k).

13th position in the private leaderboard of the corresponding Kaggle competition⁶. We attribute this relatively good score to the fact that it is obtained with the LGBM model, which was not available at the time of the competition. If the XGBh model were to be used instead, the rank would drop to the 58th position.

6.4.1 Confusion Matrix

To access the classification performance of the final LGBM model, we study the so-called confusion matrix shown in Fig. 14. True and predicted customer responses are plotted for the test sample of 36,308 customers not used for the model training (25% of the total sample). Out of 8,400 customers who responded to the offer, the model correctly predicted 2,400 (Recall = 30%). The fraction of these correctly classified positive responses in all the predicted 3,400 positive responses is 0.7 (Precision = 70%). Applying these numbers to our business problem means that the model allows us to target 30% of the customers of interest, and we should expect responses to 70% of the offers sent to customers. For completeness, the values of two other commonly used metrics are Accuracy = 80% and F1-score = 41%.

⁶<https://www.kaggle.com/c/springleaf-marketing-response/leaderboard>




Figure 15. Feature importance for the LGBM model and the reduced dataset (360 features).

6.4.2 Feature Importance

Naturally, only a limited model explainability can be preformed for a problem with anonymized information. Nevertheless, we will briefly explore the feature importance using the final model. Fig. 15 shows the number of times each feature was used in the model training. Apart from the first 25 features, the frequency decreases almost linearly with the number of features. This structure is quite different from a typical distribution for decision-tree-based ensemble models and wide datasets, characterized by a sharp excess for the most essential features followed by a long tail. Such a difference implies that the current dataset contains little redundant information, already removed using the RFE technique. The five most important features are VAR_0075, VAR_0004, VAR_0217, VAR_0322, and the engineered feature with the city density information (Sec. 4.1).

6.5 Neural-Network Based Classifier

In the final step, we build and test the performance of a model based on the neural network, using the reduced dataset with 360 features (RFE, LGMB). The model architecture comprises the input layer with 360 nodes, three hidden layers, and the one-node output layer. The so-called relu activation function is used for the hidden layers, while for the output layer, we use the sigmoid function, suitable for the binary classification task. As a part of the model hyperparameter tuning, several combinations of the number of nodes in the hidden layers are tested. We implement 360 nodes in all three layers, divide by two the number of nodes in subsequent layers, multiply and divide by two, etc. The considered architectures are listed in Tab. 7. We also test a model with four hidden layers and a model based on the full dataset, both with the number of nodes that decrease by two in subsequent layers. The number of free model parameters varies

	Model	auc_roc
0	NN3_360_360_360_1	0.7738
1	NN3_360_180_90_1	<u>0.7742</u>
2	NN3_360_720_180_1	0.7737
3	NN3_360_90_45_1	0.7728
4	NN4_360_180_90_45_1	0.7740
5	NN3_1800_940_470_1	0.7704

Table 7. NN models.




Figure 16. The loss function (left) and auc_roc score (right) as a function of the epoch number for the neural network model.

from 166,000 to 520,000 for the reduced dataset and is about 5.7 million for the full dataset.

The data are scaled using the StandardScaler method. The Adam optimizer, with the stochastic gradient descent and adjustable learning rate, and the binary_crossentropy loss function are used for the model training. The EarlyStopping callback terminates the training if the loss function does not improve after three iterations. The fits converge after only a few iterations and the training times are of the order of 1-2 minutes. Fig. 16 presents the loss function and the auc_roc score as a function of the epoch number for one of the models, showing a good correspondence between the two (the minimum loss corresponds to the maximum score). The best auc_roc score of 0.7742 is obtained for the NN3_360_180_90_1 model.

The NN-based result is better than obtained for LR and compatible with the results of the RF model after the manual tune and the gradient boosting models with the default set of parameters (Tab. 5). The score obtained for the reduced dataset is higher by a few permille than the one derived for the entire dataset, consistent with a similar observation for LGBM.

7. Summary

This study aimed to predict customer adoption using binary classification methods and a wide dataset with 2,000 anonymized features. We explored the performance of several decision-tree-based ensemble models, such as Random Forest, XGBoost, and LightGBM, as well as Logistic Regression and neural-network-based models. We found that LightGBM outperforms other models for this dataset. In terms of the area-under-the-ROC-curve, the best score of 0.8 is obtained after reducing the dimensionality of the dataset to 360 features using the Recursive Feature Elimination technique. The best model allows the Springleaf company to target 30% of the customers of interest (recall) and expect positive responses for 70% of the offers sent to them (precision).

Future work should focus on improving the relatively low recall value. It should be possible with additional domain knowledge provided by Springleaf by unveiling the data features. In particular, one should better understand the structure of the integer-type features with large adjacent positive numbers that are separated from the rest of the spectra, as mentioned in Sec. 4.2.3.

GradientBoosting (GB)	XGBoost (XGB)	LightGBM (LGBM)	HistGradientBoosting (HGB)
learning_rate = 0.1	learning_rate = 0.3	learning_rate = 0.1	learning_rate = 0.1
n_estimators = 100	n_estimators = 100	n_estimators = 100 (num_iterations)	max_iter = 100
max_depth = 3	max_depth = 6	max_depth = None	max_depth = None
max_leaf_nodes = None	-	num_leaves = 31	max_leaf_nodes = 31
min_samples_split = 2	-	-	-
mn_samples_leaf = 1	min_child_weight = 1	min_child_samples = 20 (min_data_in_leaf)	min_samples_leaf = 20
-	colsample_bytree = 1.	colsample_bytree = 1. (feature_fraction)	-
max_features = 1.	colsample_bynode = 1.	feature_fraction_bynode = 1.	-
subsample = 1.	subsample = 1.	subsample = 1. (bagging_fraction)	-
-	gamma = 1.	-	-
-	reg_lambda = 0.	reg_lambda = 0. (lambda_11)	-
-	reg_alpha = 1.	reg_alpha = 0. (lambda_12)	l2_regularization = 0.
-	max_bin = 256 tree_method = hist	max_bin = 255	max_bins = 255 (≤255)
-	n_jobs = -1	n_jobs = -1	n_jobs = -1

Table 8. Default values of hyperparameters of gradient boosting models. For the XGB and LGBM models, the names (in the brackets) correspond to those present in the scikit-learn API (stand-alone library).